

## Task-D: Collinear features and their effect on linear models

In [4]:

```
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
import pandas as pd
from sklearn.metrics import accuracy_score, roc_auc_score
import numpy as np
from sklearn.datasets import load_iris
from sklearn.linear_model import SGDClassifier, SGDRegressor
from sklearn.model_selection import GridSearchCV
import seaborn as sns
import matplotlib.pyplot as plt
import random
from tqdm import tqdm
import math
import random
```

In [5]:

```
data = pd.read_csv('/content/drive/MyDrive/temp/Linear model/task_d.csv')
```

In [6]:

```
data.tail()
```

Out[6]:

	x	y	z	x*x	2*y	2*z+3*x*x	w	target
95	0.358663	-0.207835	0.928390	0.300024	-0.207835	0.382111	1.716476	1
96	0.358663	-0.417770	0.928390	0.300024	-0.417770	0.382111	1.841269	1
97	1.141771	-0.417770	0.997724	1.130131	-0.417770	1.137586	0.019860	1
98	-0.581066	-1.257507	0.096375	-0.604025	-1.257507	-0.531992	0.590582	1
99	0.358663	-0.627704	0.859055	0.300024	-0.627704	0.373740	-0.045364	1

In [11]:


```
X = data.drop(['target'], axis=1).values
Y = data['target'].values
```


In [8]:

```
print(data.corr())
dataplot = sns.heatmap(data.corr(), cmap="YlGnBu", annot=True)
```

	x	y	z	...	2*z+3*x*x	w	target
x	1.000000	-0.205926	0.812458	...	0.996252	0.583277	0.728290
y	-0.205926	1.000000	-0.602663	...	-0.261123	-0.401790	-0.690684
z	0.812458	-0.602663	1.000000	...	0.847163	0.674486	0.969990
x*x	0.997947	-0.209289	0.807137	...	0.997457	0.583803	0.719570
2*y	-0.205926	1.000000	-0.602663	...	-0.261123	-0.401790	-0.690684
2*z+3*x*x	0.996252	-0.261123	0.847163	...	1.000000	0.606860	0.764729
w	0.583277	-0.401790	0.674486	...	0.606860	1.000000	0.641750
target	0.728290	-0.690684	0.969990	...	0.764729	0.641750	1.000000

[8 rows x 8 columns]





## Doing perturbation test to check the presence of collinearity

Indented block

### Task: 1 Logistic Regression

- Finding the Correlation between the features
  - check the correlation between the features
  - plot heat map of correlation matrix using seaborn heatmap
- Finding the best model for the given data
  - Train Logistic regression on data(X,Y) that we have created in the above cell
  - Find the best hyper parameter alpha with hyper parameter tuning using k-fold cross validation (grid search CV or random search CV make sure you choose the alpha in log space)
  - Create a new Logistic regression with the best alpha (search for how to get the best hyper parameter value), name the best model as 'best\_model'
- Getting the weights with the original data
  - train the 'best\_model' with X, Y
  - Check the accuracy of the model 'best\_model\_accuracy'
  - Get the weights W using best\_model.coef\_
- Modifying original data
  - Add a noise (order of  $10^{-2}$ ) to each element of X and get the new data set X' ( $X' = X + e$ )
  - Train the same 'best\_model' with data (X', Y)
  - Check the accuracy of the model 'best\_model\_accuracy\_edited'
  - Get the weights W' using best\_model.coef\_
- Checking deviations in metric and weights
  - find the difference between 'best\_model\_accuracy\_edited' and 'best\_model\_accuracy'
  - find the absolute change between each value of W and W'  $\Rightarrow |W - W'|$
  - print the top 4 features which have higher % change in weights compare to the other feature

### Task: 2 Linear SVM

- Do the same steps (2, 3, 4, 5) we have done in the above task 1.

**Do write the observations based on the results you get from the deviations of weights in both Logistic Regression and linear SVM**

In [16]:

```
def RandomSearchCV(x,y, params, folds):
    trainscores = [] #list to store train score for each value of k
    testscores = [] #list to store test score for each value of k
    for k in tqdm(params):

        trainscores_folds = [] #list to store train score for each fold
        testscores_folds = [] #list to store test score for each fold
        for j in range(0, folds):
            # check this out: https://stackoverflow.com/a/9755548/4084039
            #train_indices = randomly_select_60_percent_indices_in_range_from_1_to_len(x_train)
            #test_indices = list(set(list(range(1, len(x_train)))) - set(train_indices))
            list1=[*range(1,len(x)+1)]# list of all indices in x
            list2=np.array(list1)# converting it to np array
            splited_array=np.array_split(list2,folds) #splitting x to in f folds
            test_indice =splited_array[j]
            train_indice=list(set(list(range(1,len(x)+1))) - set(test_indice))#getting index fo
r all points of train
            test_indices=[x -1 for x in test_indice ] #subtracting each index by one as origina
l index of x start with 0
            train_indices=[x -1 for x in train_indice ]#subtracting each index by one as origina
l index of x start with 0
            # selecting the data points based on the train_indices and test_indices
            X_train = x[train_indices]
            Y_train = y[train_indices]
            X_test = x[test_indices]
            Y_test = y[test_indices]
            c = k
            classifier=SGDClassifier(loss='log',penalty='l2', alpha=c, l1_ratio=0.15,
fit_intercept=True, max_iter=1000, tol=0.001, shuffle=Tr
ue,
verbose=0, epsilon=0.1, n_jobs=None, random_state=None,
learning_rate='optimal', eta0=0.0, power_t=0.5, early_st
opping=False,
validation_fraction=0.1, n_iter_no_change=5, class_weigh
t=None,
warm_start=False, average=False)
            classifier.fit(X_train,Y_train)
            Y_predicted = classifier.predict(X_test)#predict y of X_test

            testscores_folds.append(accuracy_score(Y_test, Y_predicted))
            Y_predicted = classifier.predict(X_train)#predict y of X_train
            trainscores_folds.append(accuracy_score(Y_train, Y_predicted))
            trainscores.append(np.mean(np.array(trainscores_folds)))
            testscores.append(np.mean(np.array(testscores_folds)))
        return trainscores, testscores
#model = LogisticRegression() #define classifier
#list3=[*range(1,1000)]
#param=random.sample(list3,10) #generating 10 random number betwn 1 to 50
param=[0.000001,0.00001,0.0001,0.001,0.1,1,10,100,1000]
param1=[]
for i in param:
    tmp=math.log(i)
    param1.append(tmp)

params=[param.sort()]

folds = 4
trainscores, testscores = RandomSearchCV(X, Y, param, folds)

plt.plot(param1, trainscores, 'o-', label='train cruve')
plt.plot(param1, testscores, 'o-', label='test cruve')
plt.title('Hyper-parameter VS accuracy plot')
plt.legend()
plt.show()

bestmodel=SGDClassifier(loss='log',penalty='l2', alpha=10, l1_ratio=0.15,
fit_intercept=True, max_iter=1000, tol=0.001, shuffle=Tr
ue,
```

```

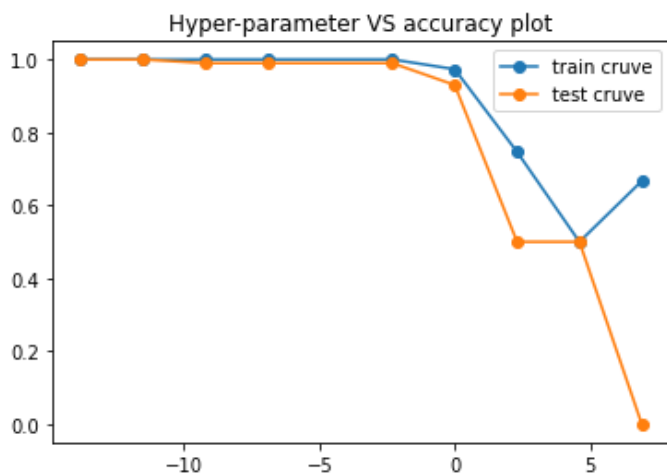
verbose=0, epsilon=0.1, n_jobs=None, random_state=None,
learning_rate='optimal', eta0=0.0, power_t=0.5, early_st
opping=False,
validation_fraction=0.1, n_iter_no_change=5, class_weigh
t=None,
warm_start=False, average=False)
bestmodel.fit(X,Y)
Y_predicted = bestmodel.predict(X) #predict y of X_test
best_model_accuracy=accuracy_score(Y, Y_predicted)
bestmodelcoef=bestmodel.coef_
print("best_model_accuracy=",best_model_accuracy)
print("feature weights=",bestmodelcoef)

X_=X+random.randrange(1,10)*0.01
bestmodel.fit(X_,Y)
Y_predicted = bestmodel.predict(X) #predict y of X_test
best_model_accuracy_edited=accuracy_score(Y, Y_predicted)
bestmodelcoef_edited=bestmodel.coef_
print("best_model_accuracy_edited=",best_model_accuracy_edited)
print("feature weights_edited=",bestmodelcoef_edited)

differene_between_accuracy=best_model_accuracy_edited-best_model_accuracy
print("differene_between_accuracy=",differene_between_accuracy)
list3=[abs(i-j) for i,j in zip(bestmodelcoef,bestmodelcoef_edited)]
print(list3)
list4=list3[0]
list4=np.array(list4)
print("abs_change_in_weights=",list4)
print(np.argsort(list4))
list_sorted=np.argsort(list4)[::-1]
print("sorted index according to abs weights=",list_sorted)
col_list=data.columns
print(col_list)
top4_feature_change=[]
for i in list_sorted[0:4]:
    top4_feature_change.append(col_list[i])
print("top4_feature_change",top4_feature_change)

```

100%|██████████| 9/9 [00:00<00:00, 196.44it/s]



```

best_model_accuracy= 0.52
feature weights= [[ 0.03244351 -0.03153831  0.0440694   0.03200811 -0.03153831  0.0341281
 9
 0.02862859]]
best_model_accuracy_edited= 0.51
feature weights_edited= [[ 0.03187996 -0.0322185   0.04350955  0.03144423 -0.0322185   0.
03356671
 0.02805848]]
differene_between_accuracy= -0.010000000000000009
[array([0.00056355, 0.00068019, 0.00055985, 0.00056389, 0.00068019,
        0.00056148, 0.00057011])]
abs_change_in_weights= [0.00056355 0.00068019 0.00055985 0.00056389 0.00068019 0.00056148
 0.00057011]
[2 5 0 3 6 1 4]
sorted index according to abs weights= [4 1 6 3 0 5 2]
Index(['x', 'y', 'z', 'x*x', '2*y', '2*z+3*x*x', 'w', 'target'], dtype='object')

```

```
top4_feature_change ['2*y', 'y', 'w', 'x*x']
```

```
In [20]:
```

```
def RandomSearchCV(x,y, params, folds):
    trainscores = [] #list to store train score for each value of k
    testscores = [] #list to store test score for each value of k
    for k in tqdm(params):

        trainscores_folds = [] #list to store train score for each fold
        testscores_folds = [] #list to store test score for each fold
        for j in range(0, folds):
            # check this out: https://stackoverflow.com/a/9755548/4084039
            #train_indices = randomly_select_60_percent_indices_in_range_from_1_to_len(x_train)
            #test_indices = list(set(list(range(1, len(x_train)))) - set(train_indices))
            list1=[*range(1,len(x)+1)]# list of all indices in x
            list2=np.array(list1)# converting it to np array
            splited_array=np.array_split(list2,folds) #splitting x to in f folds
            test_indice =splited_array[j]
            train_indice=list(set(list(range(1,len(x)+1))) - set(test_indice))#getting index fo
r all points of train
            test_indices=[x -1 for x in test_indice ] #subtracting each index by one as origina
l index of x start with 0
            train_indices=[x -1 for x in train_indice ]#subtracting each index by one as origina
l index of x start with 0
            # selecting the data points based on the train_indices and test_indices
            X_train = x[train_indices]
            Y_train = y[train_indices]
            X_test = x[test_indices]
            Y_test = y[test_indices]
            c = k
            classifier=SGDClassifier(loss='hinge',penalty='l2', alpha=c, l1_ratio=0.15,
                                     fit_intercept=True, max_iter=1000, tol=0.001, shuffle=Tr
ue,
                                     verbose=0, epsilon=0.1, n_jobs=None, random_state=None,
                                     learning_rate='optimal', eta0=0.0, power_t=0.5, early_st
opping=False,
                                     validation_fraction=0.1, n_iter_no_change=5, class_weigh
t=None,
                                     warm_start=False, average=False)
            classifier.fit(X_train,Y_train)
            Y_predicted = classifier.predict(X_test)#predict y of X_test

            testscores_folds.append(accuracy_score(Y_test, Y_predicted))
            Y_predicted = classifier.predict(X_train)#predict y of X_train
            trainscores_folds.append(accuracy_score(Y_train, Y_predicted))
            trainscores.append(np.mean(np.array(trainscores_folds)))
            testscores.append(np.mean(np.array(testscores_folds)))
        return trainscores, testscores
#model = LogisticRegression() #define classifier
#list3=[*range(1,1000)]
#param=random.sample(list3,10) #generating 10 random number betwnn 1 to 50
param=[0.000001,0.00001,0.0001,0.001,0.1,1,10,100,1000]
param1=[]
for i in param:
    tmp=math.log(i)
    param1.append(tmp)

params=[param.sort()]

folds = 4
trainscores, testscores = RandomSearchCV(X, Y, param, folds)

plt.plot(param1, trainscores, 'o-', label='train cruve')
plt.plot(param1, testscores, 'o-', label='test cruve')
plt.title('Hyper-parameter VS accuracy plot')
plt.legend()
plt.show()

bestmodel=SGDClassifier(loss='hinge',penalty='l2', alpha=10, l1_ratio=0.15,
                        fit_intercept=True, max_iter=1000, tol=0.001, shuffle=Tr
```

```


ue,
verbose=0, epsilon=0.1, n_jobs=None, random_state=None,
learning_rate='optimal', eta0=0.0, power_t=0.5, early_st
opping=False,
validation_fraction=0.1, n_iter_no_change=5, class_weigh
t=None,
warm_start=False, average=False)
bestmodel.fit(X,Y)
Y_predicted = bestmodel.predict(X) #predict y of X_test
best_model_accuracy=accuracy_score(Y, Y_predicted)
bestmodelcoef=bestmodel.coef_
print("best_model_accuracy=",best_model_accuracy)
print("feature weights=",bestmodelcoef)

X_=X+random.randrange(1,10)*0.01
bestmodel.fit(X_,Y)
Y_predicted = bestmodel.predict(X) #predict y of X_test
best_model_accuracy_edited=accuracy_score(Y, Y_predicted)
bestmodelcoef_edited=bestmodel.coef_
print("best_model_accuracy_edited=",best_model_accuracy_edited)
print("feature weights_edited=",bestmodelcoef_edited)

differene_between_accuracy=best_model_accuracy_edited-best_model_accuracy
print("differene_between_accuracy=",differene_between_accuracy)
list3=[abs(i-j) for i,j in zip(bestmodelcoef,bestmodelcoef_edited)]
print(list3)
list4=list3[0]
list4=np.array(list4)
print("abs_change_in_weights=",list4)
print(np.argsort(list4))
list_sorted=np.argsort(list4[::-1])
print("sorted index according to abs weights=",list_sorted)
col_list=data.columns
print(col_list)
top4_feature_change=[]
for i in list_sorted[0:4]:
    top4_feature_change.append(col_list[i])
print("top4_feature_change",top4_feature_change)

```

100%|██████████| 9/9 [00:00<00:00, 220.21it/s]



```

best_model_accuracy= 0.63
feature weights= [[ 0.06749926 -0.06755077  0.09295556  0.06704618 -0.06755077  0.0715649
1
0.05872509]]
best_model_accuracy_edited= 0.61
feature weights_edited= [[ 0.06717658 -0.0661877  0.0924363  0.06671768 -0.0661877  0.
07120013
0.05817939]]
differene_between_accuracy= -0.020000000000000018
[array([0.00032268, 0.00136307, 0.00051926, 0.0003285 , 0.00136307,
0.00036478, 0.0005457 ])]
abs_change_in_weights= [0.00032268 0.00136307 0.00051926 0.0003285 0.00136307 0.00036478
0.0005457 ]
[0 3 5 2 6 1 4]
sorted index according to abs weights= [4 1 6 2 5 3 0]

```

```
Index(['x', 'y', 'z', 'x*x', '2*y', '2*z+3*x*x', 'w', 'target'], dtype='object')
top4_feature_change ['2*y', 'y', 'w', 'z']
```

## **OBSERVATION**

**1.In linear regression model after adding error to the features,we do not see much change in weights after pertubation.**

**2.In LR, absolute changes in weights are approx  $10^{-3}$**

**3.accuracy of the model also decreased a little bit .**

**4.In SVM also feature wieghts do not change much after pertubation, so features are not collinear .**

**5.accuracy score difference is also less.**

**6.we can conclude that when feature weights drastically after pertubation ,then the feature whose weights changed are collinear feature.**